



ABERYSTWYTH UNIVERSITY

COMPUTER SCIENCE AND STATISTICS (GG34)

CS396: MINOR PROJECT

Application of Machine Learning Techniques to Next Generation Sequencing Quality Control

Author:
Sam Nicholls msn

Supervisor:
Dr. Amanda Clare afc

Draft
April 2, 2014

Declaration

I certify that except where indicated, all material in this thesis is the result of my own investigation and references used in preparation of the text have been cited. The work has not previously been submitted as part of any other assessed module, or submitted for any other degree or diploma.

Sam Nicholls

2014

Abstract

Over the past few years advances in genetic sequencing hardware have introduced the concept of massively parallel DNA sequencing; allowing potentially billions of chemical reactions to occur simultaneously, reducing both time and cost required to perform genetic analysis[3]. However, these "next-generation" processes are complex and open to error[2], thus quality control is an essential step to assure confidence in any downstream analyses performed.

During sample sequencing a large number of quality control metrics are generated to determine the quality of the reads from the sequencing hardware itself. At the Wellcome Trust Sanger Institute, the automated QC system currently relies on hard thresholds to make such quality control decisions with individual hard-coded values on particular metrics determining whether a lane has reached a level that requires a warning, or has exceeded the threshold and failed entirely. Whilst this does catch most of the very poor quality lanes, a large number of lanes are flagged for manual inspection at the warning level; a time consuming task which invites inefficiency and error.

In practise most of these manual decisions are based on inspecting a range of diagnostic plots which suggests that a machine learning classifier could potentially be trained on the combinations of quality control statistics available to make these conclusions without the need for much human intervention.

Contents

Contents	iii
1 Introduction	1
1.1 Project Aims	1
1.1.1 Analysis of Current System	1
1.1.2 Identification of Properties that affect Downstream Analysis	2
2 Analysis of Current System	3
2.1 Introduction	3
2.1.1 Data Collection and Format	3
Appendix A samtools stats example output	5

Chapter 1

Introduction

Over the past few years advances in genetic sequencing hardware have introduced the concept of massively parallel DNA sequencing; allowing potentially billions of chemical reactions to occur simultaneously, reducing both time and cost required to perform genetic analysis[3]. However, these "next-generation" processes are complex and open to error[2], thus quality control is an essential step to assure confidence in any downstream analyses performed.

1.1 Project Aims

The project consists of two sub-projects;

- Analysis of a current quality control system in place
- Identification of quantifiable sample properties that affect downstream analysis

1.1.1 Analysis of Current System

With the support of the Wellcome Trust Sanger Institute in Cambridge, this project works with the Human Genetics Informatics team to investigate **auto_qc**, the institute's current automated quality control tool.

During genetic sequencing a large number of metrics are generated to determine the quality of the data read from the sequencing hardware itself. As part of the current vertebrate sequencing pipeline[1] at the institute, **auto_qc** is responsible for applying quality control to samples within the pipeline by comparing a modest subset of these metrics to simple hard-coded hard thresholds; determining whether a particular sample has reached a level that requires a warning, or has exceeded the threshold and failed entirely. Whilst this does catch most of the very poor quality outputs, a large number of samples are flagged for manual inspection at the warning level; a time consuming task which invites both inefficiency and error.

In practise most of these manual decisions are based on inspecting a range of diagnostic plots which suggests that a machine learning classifier could potentially be trained on the combinations of quality control statistics available to make these conclusions without the need for much human intervention.

The first part of the project aims to apply machine learning techniques to replicate the current **auto_qc** rule set by training a decision tree classifier on a large set of these quality metrics. The idea is to investigate whether these simple threshold based rules can be recovered from such data, or whether a new classifier would produce different rules entirely. During this analysis it is hoped the classifier may be able to identify currently unused quality metrics that improve labelling accuracy. An investigation on the possibility of aggregating or otherwise reducing the dimensions of some of the more detailed quality statistics to create new parameters will also be conducted.

The goal is to improve efficiency of quality control classification, whether by improving accuracy of pass and fail predictions over the current system or merely being able to provide additional information to a lab technician inspecting samples labelled with a warning to reduce arbitrary decisions.

1.1.2 Identification of Properties that affect Downstream Analysis

The other half of this project is motivated by the question "What *is* good and bad in terms of quality?"

To be able to classify samples as a pass or a fail with understanding, we need an idea of what actually constitutes a good or bad quality sample and must look at the effects quality has on analysis performed downstream from sequencing. An example of such is **variant calling** — the process of identifying differences between a DNA sample (such as your own) and a known reference sequence.

Given two high quality data sources where DNA sequences from individuals were identified in two different ways (one of which being next-generation sequencing) it would be possible to measure the difference between each corresponding pair. Using this, we could investigate the effect of leaving out part of the next-generation sample during the variant calling process. If we were to leave a part of a sample out of the variant calling pipeline would the variants found be more (or less) accurate than if it had been included? Would they agree more (or less) with the variants called after using the non next-generation sequencing method?

Having identified such sub-samples, can quality control metrics from the previous part be found in common? If so, such parameters would identify "good" or "bad" samples straight out of the machine! Samples that exhibit these quality variables will go on to improve or detriment analysis.

Chapter 2

Analysis of Current System

2.1 Introduction

2.1.1 Data Collection and Format

As part of the project I've been granted access to significant data sets at the Sanger Institute, unlocking quality control data for two of the largest studies currently undergoing analysis. A wide array of quality metrics are available for each and every lanelet that forms part of either of the two studies; totalling 13,455 files.

The files are created by **samtools stats** — part of a collection of widely used open-source utilities for post processing and manipulation of large alignments such as those produced by next-generation sequencers that are released under the umbrella name of "SAMtools" (Sequence Alignment and Map Tools). **samtools stats** collects statistics from sequence data files and produces key-value summary numbers as well as more complex tab delimited dataframes tabulating several metrics over time.

The output of **samtools stats** is then parsed by an in-house tool called **bamcheckr**, named so as **samtools stats** was once known as **bamcheck** and the tool is written in R. **bamcheckr** supplements the summary numbers section of the **samtools stats** output with additional metrics that are later used by **auto_qc** for classification. This process does not change the file other than adding a few additional key-value pairs in the summary numbers section. An truncated example of a "bamcheckr'd" file can be found in Appendix A.

References

- [1] vr-pipe, a generic pipeline system [Github]. [Online]. Available: <https://github.com/wtsi-hgi/vr-pipe/>
- [2] M. Kircher, U. Stenzel and J. Kelso, "Improved base calling for the Illumina Genome Analyzer using machine learning strategies," *Genome Biology*, vol. 10, no. 8, p. R83, 2009.
Useful introduction to relevant Illumina hardware and the errors that can occur during sequencing.
- [3] T. Strachan and A. Read, *Human Molecular Genetics*, 4th ed. Garland Science, 2011, pp. 214–254.
A concise introduction to the processes involved in massively parallel DNA sequencing.

Appendix A

samtools stats example output

```
# Summary Numbers. Use 'grep ^SN | cut -f 2-' to extract this part.
SN      raw total sequences:      41400090
SN      filtered sequences:       0
SN      sequences:                41400090
SN      is paired:                1
SN      is sorted:                1
SN      1st fragments:            20700045
SN      last fragments:           20700045
SN      reads mapped:             41291484
SN      reads unmapped:           108606
SN      reads unpaired:           60000
SN      reads paired:             41231484
SN      reads duplicated:         5756822
SN      reads MQ0:                1038644
SN      reads QC failed:         0
SN      non-primary alignments:    0
SN      total length:             3105006750
SN      bases mapped:            3096861300
SN      bases mapped (cigar):      3090885143
SN      bases trimmed:           0
SN      bases duplicated:         431761650
SN      mismatches:              9107833
SN      error rate:              0.002946675
SN      average length:          75
SN      maximum length:          75
SN      average quality:          36
SN      insert size average:      178.7
SN      insert size standard deviation: 44.1
SN      inward oriented pairs:    20577242
SN      outward oriented pairs:   3140
SN      pairs with other orientation: 3711
SN      pairs on different chromosomes: 31535
SN      fwd percent insertions above baseline: 1.43135383851191
SN      fwd percent insertions below baseline: 0.686265539012562
SN      fwd percent deletions above baseline: 1.38326380878871
SN      fwd percent deletions below baseline: 0.44923551909251
SN      rev percent insertions above baseline: 1.08264446659241
SN      rev percent insertions below baseline: 0.457290262062496
SN      rev percent deletions above baseline: 1.15931214598243
SN      rev percent deletions below baseline: 0.413119424753248
SN      contiguous cycle dropoff count: 36
SN      fwd.percent.insertions.above.baseline: 1.43135383851191
```

```

SN      fwd.percent.insertions.below.baseline:      0.686265539012562
SN      fwd.percent.deletions.above.baseline:        1.38326380878871
SN      fwd.percent.deletions.below.baseline:        0.44923551909251
SN      rev.percent.insertions.above.baseline:        1.08264446659241
SN      rev.percent.insertions.below.baseline:        0.457290262062496
SN      rev.percent.deletions.above.baseline:        1.15931214598243
SN      rev.percent.deletions.below.baseline:        0.413119424753248
SN      quality.dropoff.fwd.high.iqr.start.read.cycle:      0
SN      quality.dropoff.fwd.high.iqr.end.read.cycle:        0
SN      quality.dropoff.fwd.high.iqr.max.contiguous.read.cycles:      0
SN      quality.dropoff.fwd.mean.runmed.decline.start.read.cycle:      20
SN      quality.dropoff.fwd.mean.runmed.decline.end.read.cycle:      51
SN      quality.dropoff.fwd.mean.runmed.decline.max.contiguous.read.cycles:      32
SN      quality.dropoff.fwd.mean.runmed.decline.high.value:      36.9775883578997
SN      quality.dropoff.fwd.mean.runmed.decline.low.value:      36.301749247405
SN      quality.dropoff.rev.high.iqr.start.read.cycle:      0
SN      quality.dropoff.rev.high.iqr.end.read.cycle:        0
SN      quality.dropoff.rev.high.iqr.max.contiguous.read.cycles:      0
SN      quality.dropoff.rev.mean.runmed.decline.start.read.cycle:      18
SN      quality.dropoff.rev.mean.runmed.decline.end.read.cycle:      56
SN      quality.dropoff.rev.mean.runmed.decline.max.contiguous.read.cycles:      39
SN      quality.dropoff.rev.mean.runmed.decline.high.value:      36.1517621338504
SN      quality.dropoff.rev.mean.runmed.decline.low.value:      35.3152133727245
SN      quality.dropoff.high.iqr.threshold:            10
SN      quality.dropoff.runmed.k:                      25
SN      quality.dropoff.ignore.edge.cycles:            3
SN      A.percent.mean.above.baseline:                0.0991164444444441
SN      C.percent.mean.above.baseline:                0.1273795555555556
SN      G.percent.mean.above.baseline:                0.0603679999999997
SN      T.percent.mean.above.baseline:                0.0868000000000005
SN      A.percent.mean.below.baseline:                0.09911644444444451
SN      C.percent.mean.below.baseline:                0.1273795555555555
SN      G.percent.mean.below.baseline:                0.0603680000000002
SN      T.percent.mean.below.baseline:                0.0867999999999993
SN      A.percent.max.above.baseline:                 0.6017333333333332
SN      C.percent.max.above.baseline:                 0.3942666666666667
SN      G.percent.max.above.baseline:                 0.2956
SN      T.percent.max.above.baseline:                 0.7680000000000001
SN      A.percent.max.below.baseline:                 0.3182666666666666
SN      C.percent.max.below.baseline:                 0.8257333333333332
SN      G.percent.max.below.baseline:                 0.5544000000000001
SN      T.percent.max.below.baseline:                 0.2519999999999999
SN      A.percent.max.baseline.deviation:             0.6017333333333332
SN      C.percent.max.baseline.deviation:             0.8257333333333332
SN      G.percent.max.baseline.deviation:             0.5544000000000001
SN      T.percent.max.baseline.deviation:             0.7680000000000001
SN      A.percent.total.mean.baseline.deviation:      0.1982328888888889
SN      C.percent.total.mean.baseline.deviation:      0.2547591111111111
SN      G.percent.total.mean.baseline.deviation:      0.120736
SN      T.percent.total.mean.baseline.deviation:      0.1736

```

First Fragment Qualities. Use 'grep ^FFQ | cut -f 2-' to extract this part.

Columns correspond to qualities and rows to cycles. First column is the cycle number.

FFQ	1	8968	3619	9863	747	5094	0	6642	1609	4673	4208	20
FFQ	2	21676	0	0	0	0	0	0	43	1885	0	0
FFQ	3	7	0	177	0	0	0	0	0	0	0	0
FFQ	4	0	0	0	65	0	0	0	272	0	0	14277
FFQ	5	1917	173	1249	0	1890	0	0	0	10874	0	0
[...]												
FFQ	72	4098	0	0	4806	0	0	0	65507	0	0	0
FFQ	73	3894	2	0	0	0	4931	53483	0	0	0	0
FFQ	74	3697	39	0	919	4933	0	0	56866	1524	0	0
FFQ	75	4542	0	0	0	0	0	4634	77822	0	0	0
FFQ	76	0	0	0	0	0	0	0	0	0	0	0

```

# Last Fragment Qualities. Use 'grep ^LFQ | cut -f 2-' to extract this part.
# Columns correspond to qualities and rows to cycles. First column is the cycle number.
LFQ 1 8869 0 0 0 0 0 63 0 0 1156 616 173
LFQ 2 3300 0 0 0 0 0 0 0 0 0 389 0 0
LFQ 3 6816 0 0 0 573 0 83 0 7011 1171 107134 0
LFQ 4 5492 0 0 13 0 66 730 708 8134 2422 84052
LFQ 5 3512 0 0 0 1023 185 0 8653 1995 0 115559
[...]
LFQ 72 5135 166 0 0 2872 13643 0 59649 4249 11351 346
LFQ 73 6025 229 0 86 1042 13417 0 66093 3741 8151 354
LFQ 74 5980 3 91 0 0 1340 9696 72939 4924 304090
LFQ 75 4314 0 0 168 0 848 8591 0 70358 3827 352180
LFQ 76 0 0 0 0 0 0 0 0 0 0 0 0
# Mismatches per cycle and quality. Use 'grep ^MPC | cut -f 2-' to extract this part.
# Columns correspond to qualities, rows to cycles. First column is the cycle number, second
# is the number of N's and the rest is the number of mismatches
MPC 1 14078 0 2594 6777 416 1919 0 2222 352 987 537
MPC 2 21407 0 0 0 0 0 0 0 5 223 19 0
MPC 3 3205 0 0 37 0 43 0 12 0 691 71 6984
MPC 4 1774 0 0 0 2 29 4 65 73 863 192 6749
MPC 5 1913 0 94 885 0 969 23 0 959 213 1203 84
[...]
MPC 72 361 0 13 0 573 276 934 0 16376 426 1066
MPC 73 1005 0 11 0 4 79 777 539 15025 363 699
MPC 74 779 0 3 0 131 440 0 93 7485 6589 387 24
MPC 75 136 0 0 0 3 0 47 704 9302 5886 260 2650
MPC 76 0 0 0 0 0 0 0 0 0 0 0 0
# GC Content of first fragments. Use 'grep ^GCF | cut -f 2-' to extract this part.
GCF 0.5 56
GCF 1.76 60
GCF 3.02 126
GCF 4.27 212
GCF 5.78 347
[...]
GCF 93.72 378
GCF 95.23 186
GCF 96.48 87
GCF 97.74 55
GCF 99.25 17
# GC Content of last fragments. Use 'grep ^GCL | cut -f 2-' to extract this part.
GCL 0.5 118
GCL 1.76 175
GCL 3.02 230
GCL 4.27 354
GCL 5.78 525
[...]
GCL 93.72 613
GCL 95.23 430
GCL 96.48 274
GCL 97.74 185
GCL 99.25 110
# ACGT content per cycle. Use 'grep ^GCC | cut -f 2-' to extract this part. The columns are: cycle, and A,C,G,T counts [%]
GCC 1 26.93 23.09 22.77 27.2
GCC 2 26.78 23.24 22.97 27.02
GCC 3 26.46 23.59 23.3 26.66
GCC 4 26.29 23.79 23.45 26.46
GCC 5 26.47 23.61 23.3 26.62
[...]
GCC 70 26.09 24.26 23.45 26.2
GCC 71 26.07 24.25 23.46 26.22
GCC 72 26.04 24.27 23.49 26.2
GCC 73 26.07 24.25 23.47 26.22
GCC 74 26.08 24.24 23.45 26.23

```

```

GCC      75      26.01      24.31      23.51      26.18
# Insert sizes. Use 'grep ^IS | cut -f 2-' to extract this part. The columns are: pairs total, inward oriented pairs, outward oriented pairs
IS      0      10      0      1      9
IS      1      3      0      3      0
IS      2      4      0      4      0
IS      3      5      0      5      0
IS      4      2      0      2      0
IS      5      3      0      3      0
[...]
IS      110     33952     33952      0      0
IS      111     38433     38433      0      0
IS      112     43373     43370      0      3
IS      113     48160     48159      0      1
IS      114     53175     53171      0      4
IS      115     59504     59502      0      2
IS      116     64668     64668      0      0
IS      117     71107     71105      0      2
IS      118     77157     77156      0      1
IS      119     84044     84044      0      0
IS      120     90116     90110      3      3
[...]
IS      327     6546      6546      0      0
IS      328     6483      6483      0      0
IS      329     6201      6201      0      0
IS      330     6228      6228      0      0
IS      331     5852      5852      0      0
# Read lengths. Use 'grep ^RL | cut -f 2-' to extract this part. The columns are: read length, count
RL      75      41400090
# Indel distribution. Use 'grep ^ID | cut -f 2-' to extract this part. The columns are: length, number of insertions, number of deletions
ID      1      128650     183418
ID      2      26409      39770
ID      3      10213     16046
ID      4      7756      11444
ID      5      1746      3455
[...]
ID      35      0      8
ID      36      0      1
ID      37      0      1
ID      38      0      1
ID      40      0      2
# Indels per cycle. Use 'grep ^IC | cut -f 2-' to extract this part. The columns are: cycle, number of insertions (fwd), .. (rev) , number of deletions
IC      1      0      0      105      97
IC      2      24      15      150      179
IC      3      129      138      441      509
IC      4      253      310      623      829
IC      5      557      724      786      1164
[...]
IC      70      571      710      638      761
IC      71      350      428      309      434
IC      72      154      150      38      45
IC      73      60      61      15      23
IC      74      20      19      11      12
# Coverage distribution. Use 'grep ^COV | cut -f 2-' to extract this part.
COV      [1-1]      1      332980694
COV      [2-2]      2      105004580
COV      [3-3]      3      29112182
COV      [4-4]      4      13415014
COV      [5-5]      5      6716815
[...]
COV      [996-996]      996      2
COV      [997-997]      997      2
COV      [998-998]      998      2
COV      [1000-1000]      1000      4

```

```

COV      [1000<]      1000      116
# GC-depth. Use 'grep ^GCD | cut -f 2-' to extract this part. The columns are: GC%, unique sequence percentiles, 10th, 25th, 50th, 75th and 90th
GCD      0          0.001      0          0          0          0          0
GCD      0.4        0.002      0.101      0.101      0.101      0.101      0.101
GCD      19         0.003      0.049      0.049      0.049      0.049      0.049
GCD      20         0.004      0.06       0.06       0.06       0.06       0.06
GCD      21         0.004      0.045      0.045      0.045      0.045      0.045
[...]
GCD      66         99.99      0.244      2.693      6.746      11.794     15.885
GCD      67         99.994     1.279      1.279      4.305      9.667      11.483
GCD      68         99.997     4.148      4.148      4.463      5.741      7.354
GCD      69         99.999     0.499      0.499      0.499      1.935      1.935
GCD      72         100       0.476      0.476      0.476      1.219      1.219

```